



# UE Learning from User-generated Data S2024 MRS Challenge

# Agenda

- Evaluate the course
- MRS challenge
- Implementation ideas, guidelines



# The Challenge

- Recommend songs to users of Last FM
- Goal: maximize overall nDCG

## Data: LFM2B

- Bigger (~8 times more items) LFM2b sample
- Contents:
  - lfm-challenge.user id, country, age, gender, registration date
  - Ifm-challenge.item id, artist, track, country of origin
  - Ifm-challenge.inter\_train user id, item id, listening events
  - Ifm-challenge.inter\_test user id, item id, listening events
  - Ifm-challenge.musicnn item features
  - test\_indices.txt ids of test users

## Submission – two files

 rec\_<matr.num.>\_<name>.tsv - the file with your recommendations, where each line corresponds to one target user:

```
<user_id>\t<item_id>,<item_id>,...
```

10 not seen items for each user (no spaces between), example:

rec_k0000007_Bond_James.tsv	
0 7	0,1,2,3,4,5,6,7,8,9 15,1,2,3,4,0,6,8,9,804
9	4,15,10,2,3,5,7,8,9,190
•••	

report \_<matr.num.> \_<name>.txt - text file describing your approach
to the recommendation part of the task: algorithms used, training
settings, ... No more than ½ A4; Specify your name and matr. number
within as well

## Evaluation

- The results will be judged by the overall nDCG performance and the report
- The exercise is worth 10 points with up to 10 bonus points available
  - Confidently beat POP-recommender 6 points
  - ItemKNN performance 8 points
  - Standard report 2 points
  - Additional points for better nDCG, interesting and sound methodology, deep analysis of the result, high quality reports

# Timing

- Challenge publication: 5.06.24
- Deadline: 25.06.2023
- Occasional (~weekly) score board updates
- Possible Sync/QA meeting(s)

Questions so far?



# Establish experimental framework

- 1. Create additional data splits train-(val)-test (at least 3 in total)
- 2. Setup the nDCG evaluation pipeline
- 3. Produce & evaluate random recommendations
- 4. Produce & evaluate POP recommendations
- 5. Begin experiments...

# Experiments

#### For each investigated model:

- Perform hyper-parameter search (try out different embedding sizes, number of neighbors etc.)
- 2. Evaluate nDCG of each configuration of each model on all 3+ data splits. Compare the average to the two baselines
- 3. Document all experiments: motivation, hyper-parameters, results. Save results produced by costly methods

# Approaches

- 1. Try approaches from the UE (some will require optimization)
- 2. Look up hints in literature
- 3. Investigate performance of the models, find weak spots. E.g. what users get the lowest nDCG? What can be done for them?
- 4. Hybridize, use content, try different things, have fun.
- 5. Keep your experiments logical and clearly motivated.
- 6. Use recommendation frameworks at your own risk. Make sure you understand approaches you use (every report will be checked).

# Report

- 1. Keep it informative but concise
- 2. Describe your experimental setup
- 3. Explain tried approaches and motivation of moving from one to another
- 4. Report performance of all systems including the baselines
- 5. The detail should be enough to reproduce the results. Omit code and standard implementation details
- 6. Use LLMs at your own risk uninformative reports will be penalized.

#### Ideas

- 1. Souped up POP recommender
- 2. ItemKNN
- 3. Explore Matrix Factorization
- 4. Explore more approaches

#### Extended POP recommender

#### Demographic filtering:

#### Hypotheses:

- Users with similar demographics (age, country, gender) have similar tastes
- Users enjoy a certain portion of their local music

#### Filtering by preference:

- Use more sophisticated ways to cluster users
- Users of each cluster have their own distinct top-pop

#### Additional data:

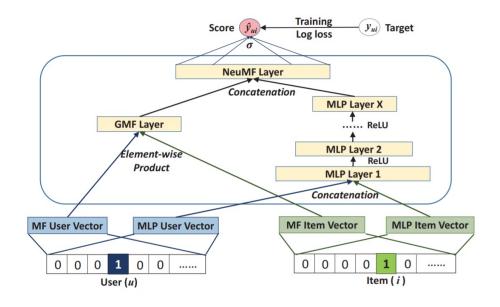
 Assign genres to tracks, make POP recommender aware of user preference genre-wise

#### Item KNN

- Solid recommender for LFM-2B
- Will require optimization for the new dataset
- Find publication introducing Item KNN
- Check optimization suggestions

## Matrix Factorization

- Optimize our implementations
- Check published approaches (SVD, ALS-MF, BPR-MF, ...)
- Neural Matrix Factorization
- Shallow autoencoders
- Content: Factorization machines, Inductive Matrix Factorization



## In a nutshell

- Start small (first try simplest approaches, get your experimental framework going)
- Tune hyper-parameters (N of neighbors, embedding sizes, ...)
- Compare your implementations to each other and baselines (e.g. POP, random recommender, different versions of your system)
- Use multiple train/test splits to make sure one system consistently outperforms the other (we only supply one split)
- Once you figured out your optimal parameters feel free to train your system on the whole data (train + test) and then produce final recommendation
- Optimize, use Google Colab or sample the data if your machine struggles with it
- Follow the submission format closely and don't forget the report!

# Thank you & good luck!